

Daimler Trucks & Buses Tech and Data Hub

Big Data Technical Challenge

Carlos Godinho

cma.godinho@gmail.com

Version 1.0, April 2019

1 Answers

Analyze 3 datasets to produce a dataset that answers the question:

Which doctors from which medical school prescribe drugs of a certain family and type?

And answer 3 proposed questions:

1 - The following table from the above 3 datasets. Please compile the data from the above 3 data sets.

Read chapter 3 - Exploratory Data Analysis for the detailed data analysis.

Check the proposal for the consolidated dataset in chapter 3.6 - Final Dataset.

A Spark application (Consolidation) is available to implement the proposal and is described in chapter 4 - Solution.

2 - In the case that the data was available as a web download (not as a downloaded file), modify the program to read from a HTTP source (S3 / Blob storage). Is it possible to process the file without local storage / staging in HDFS?

Data is available as a web download, as described in chapter 2.4 - Other Data Availability.

As available in Spark Scala doc:

<https://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.SparkContext>

Spark reads directly from HDFS, S3 or a local filesystem. Direct reading as HTTP(S) is not supported.

A specific implementation is added to Consolidation application to parse HTTP(S) input files. The implementation is available in class Manager, method webFile. The solution is working as described in chapter 4.4 - Execution. However, this solution is not recommended for large files (> 1 million records).

An alternative solution is presented in Figure 1. Data is downloaded, parsed and injected into Kafka. Spark Streaming read records from Kafka and processes data. Final data is stored in a final device.



Figure 1 – Alternative solution

Therefore, it is possible to process the file without local storage / HDFS.

3 - How can the original data be downloaded and deflated directly from spark? Write a spark job for this?

Spark application (Consolidation) execution and tests are described in chapter 0 -

Tests.

2 Datasets

2.1 Dataset Drugs

Content: The primary data source for these data is the CMS Chronic Condition Data Warehouse (CCW), a database with 100% of Medicare enrollment and final-action Part D prescription drug event (PDE) data.

Location: <https://data.cms.gov/Medicare-Part-D/Part-D-Prescriber-National-Summary-Report-Calendar/2n5w-7ghf> is the base directory for Drugs.

https://download.cms.gov/Research-Statistics-Data-and-Systems/Statistics-Trends-and-Reports/Medicare-Provider-Charge-Data/Downloads/PartD_Prescriber_PUF_Drug_Ntl_15.zip, holds the data.

Format: Microsoft Excel Open XML Spreadsheet (XLSX).

Data Dictionary: In file, tab Data Dictionary

2.2 Dataset Physicians

Content: This file contains general information about individual eligible professionals (EPs) such as demographic information and Medicare quality program participation.

Location: <https://data.medicare.gov/data/Physician-Compare> is the base directory for Physicians.

<https://data.medicare.gov/Physician-Compare/Physician-Compare-National-Downloadable-File/mj5m-pzi6> holds the data.

NOTE: that there is also a link for the register 2015 Physicians:

http://medicare.gov/download/PhysicianCompare/2015/Refresh_Data_Archive_November_2015_1.zip

this is the used link!

Format: ZIP with Comma Separator Values. Commas may exist within fields. In such case, fields are quoted.

Data Dictionary: In the page

2.3 Dataset Prescriptions

Content: The Part D Prescriber PUF is organized by National Provider Identifier (NPI) and drug name and contains information on drug utilization (claim counts and day supply) and total drug costs.

Location: http://download.cms.gov/Research-Statistics-Data-and-Systems/Statistics-Trends-and-Reports/Medicare-Provider-Charge-Data/Downloads/PartD_Prescriber_PUF_NPI_DRUG_15.zip

Format: ZIP with Tab Separator Values.

Data Dictionary: https://www.cms.gov/Research-Statistics-Data-and-Systems/Statistics-Trends-and-Reports/Medicare-Provider-Charge-Data/Downloads/Prescriber_Methods.pdf

2.4 Other Data Availability

Data.Medicare.gov data is supported by the **Socrata Open Data platform (SODA)**, which provides an API to allow software developers to access it.

The description link is presentation here: <https://data.medicare.gov/developers>

The SODA API supports formats like JSON, XML and CSV. Each dataset has a specific endpoint to access the data.

Examples on how to use the API: <https://dev.socrata.com/>

For the drugs we have as example:

<https://data.cms.gov/Medicare-Part-D/Part-D-Prescriber-National-Summary-Report-Calendar/2n5w-7ghf>

With data end points like:

<https://data.cms.gov/resource/5m6j-9i3f.json>

<https://data.cms.gov/resource/5m6j-9i3f.csv>

3 Exploratory Data Analysis

3.1 Introduction

To make an effective exploratory data analysis to the provided data, [R](#) language is used. Besides basic R the following components are required:

- RStudio for IDE support;
- R Packages (`dplyr`, `readr` and `stringr`) to increase R productivity.

The R script is available as `/code/data_analysis.R`.

3.2 Entities

3 entities are identified:

- **Drug** – lists a set of drugs with prices and properties;
- **Physician** – Defines a set of Physicians, with specialties, training, alma-mater and other personal details;
- **Prescription** – Records of drug prescriptions made by physicians during 2015.

Each entity is available as a dataset (tabular format).

3.3 Size and Keys

Dataset size and keys are presented in Table 1.

Dataset	Size	Primary KEY
Drug	3 395 x 19	Drug Name + Drug Generic Name
Physician	2 013 843 x 43	NPI is the key, although multiple updates are available There are other identifiers with no applicability for this use case
Prescription	24 524 894 x 21	As stated in data dictionary, it is an aggregation of: NPI + Drug Name + Drug Generic Name

Table 1 – Dataset size and keys

3.4 Findings

In **Drug**:

- Drugs dataset has a composite key of `Drug Name` + `Generic name`. In fact, there are only 190 repetitions of field `Drug Name`. In a real scenario a discussion with a domain expert would be recommended;
- Considering both `Drug Name` and `Generic Name`, all fields are available.

In **Physician**:

- Using file form 2015;
- Entries related to the same Physician is repeated, due to updates in its descriptions;
- The pair `<NPI, Medical school name>` can be aggregated as it never changes.

In **Prescription**:

- Field `bene_count` has ~3 million empty records. According to the data dictionary, counts fewer than 11 are suppressed and are indicated by a blank;
- Key fields are always available;
- There are 1 208 entries of empty field `nppes_provider_first_name`, which corresponds to providers registered as organizations. This is a very low rate;
- Holds references to Prescribers (by field `npi`);
- Holds references to Drugs (by field `drug_name` and `generic_name`);
- Besides individual providers, attribute NPI includes organizational providers, nursing homes, group practices, non-physician practitioners, residential facilities, ambulatory surgery centers, and other providers.

Between entities:

- There are 336 092 prescriptions pointing to unavailable Physicians. According to the data description not all the Physicians are available. Also, some Prescribers are not Physicians (although they have `npi`);
- All Drugs reported in Prescriptions are available in Drugs dataset.

3.5 Relationships

The relationships between entities are presented in Figure 2.

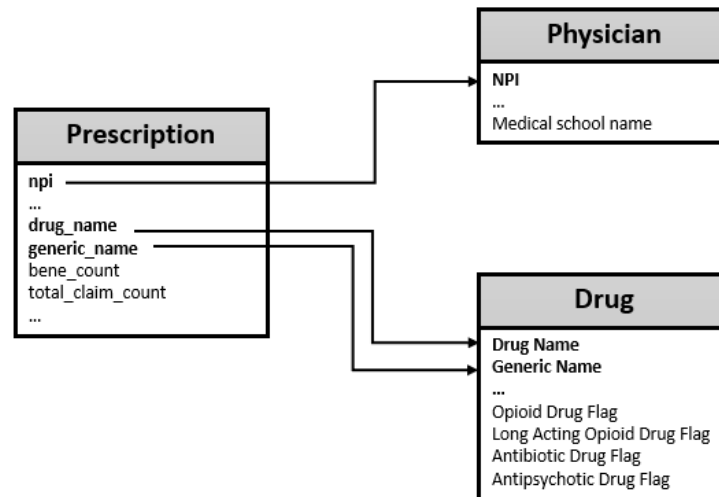


Figure 2 - Relationship between entities. Keys are represented in bold.

3.6 Final Dataset

The final dataset purposed is presented in Table 2, together with the data source location. A markdown file is available with the data dictionary as `/doc/consolidated_prescription.md`.

New Dataset		Original Entity			Notes
Name	#	Entity	Name	#	
npi	1	Physician Prescription	NPI npi	1 1	Numeric
drug_simple_name	2	Drug Prescription	Drug Name	1 8	String
drug_generic_name	3	Drug Prescription	Generic Name	2 9	String
bene_count	4	Prescription	bene_count	10	Numeric, Not always available May be useful for data scientists
total_claim_count	5	Prescription	total_claim_count	11	Numeric, May be useful for data scientists
medical_school	6	Physician	Medical school name	10	String
drug_opioid_flag	7	Drug	Opioid DF	16	Boolean Drug type descriptor
drug_long_acting_opioid_flag	8	Drug	Long Acting Opioid DF	17	Boolean Drug type descriptor
drug_antibiotic_flag	9	Drug	Antibiotic DF	18	Boolean Drug type descriptor
drug_antipsychotic_flag	10	Drug	Antipsychotic DF	19	Boolean Drug type descriptor

Table 2 - Selected fields and its origin

4 Solution

4.1 Introduction

A Spark solution validates and combines the datasets. As a result, the **Consolidation dataset** is created. To build the solution the following components are required:

- IntelliJ IDE with Scala extension;
- Scala 2.11.12
- Scala Test 3.0.5
- Spark 2.4.0
- sbt 1.2.8
- HDFS 2.9.2
- Apache HTTP server 2.2 (for simulating HTTP download)

The project is available as **Consolidation** and is available in `/code/Consolidation`.

4.2 Architecture

The following points are the result of design/architecture decisions:

- The application executes in a Spark Cluster and accesses HDFS;
- For the use underlying data structures, RDDs are used. RDDs guarantee best performance in data processing. As an alternative the SQL/Dataset/Dataframe could also be used;
- The application is available as a jar to be launched in a cluster with spark-submit;
- For flexibility, the application receives a set of arguments when launched;
- It is recommended to have input data available in HDFS, as an alternative an URL may also be provided;
- Results are provided to an output directory in HDFS;
- Input data files have fields separators, but the field separator may also be used inside a file. In this case the fields are quoted. For simplification the output data uses “,” as single separator and removes quotes from fields. In the case “,” are used in a field, it is replaced by “;”;
- Scala test is used for testing a class and to show its potentiality;
- Each application execution produces a log in the output directory (`consolidation.<date+time>.log`), detailing its activities and timestamps;
- Data is joined by **inner join** concept, if no matching is found, entries are not considered.
- sbt script is available for producing the application.

4.3 Spark Performance

The following performance improvements are available:

- **Kyro** serialization is activated to improve performance as data needs to be changed between executors;

- As Drugs have a small size compared to the other datasets, it is more efficient to transform Drugs data into a Map and use it as a lookup. Drugs data is Broadcasted into the executors as a name value Map;
- Filtering only relevant fields reduces the amount of data to be processed and shuffled in the cluster;
- As the application is working in a cluster, adding more resources (**horizontal scaling**) improves performance without the need of any changes in its internal implementation;
- Many physicians have its medical school reported as other, a possible performance improvement is to eliminate such elements during parsing (not implemented).

4.4 Execution

To run the application, use the following command:

```
$SPARK_HOME/bin/spark-submit --master spark://<spark address> consolidation_2.11-*.jar $1 $2 $3 $4 $5 $6
```

where:

- \$1 – HADOOP configuration files CSV list (`core-site.xml` and `hdfs-site.xml` are mandatory)
- \$2 – HADOOP base path
- \$3 – Drug file (must be inside HADOOP base path + `input`), if file in a web link, just provide an URL)
- \$4 – Physician file (must be inside HADOOP base path + `input`), if file in a web link, just provide an URL)
- \$5 - Prescription file (must be inside HADOOP base path + `input`), if file in a web link, just provide an URL)
- \$6 – Consolidation output file (written to HADOOP base path + `output`)

Shell scripts are available in directory `/tests/<environment>/execution`.

5 Tests

5.1 Introduction

The following test strategy is defined:

- Create a small data set to verify the consolidation business logic. Sample files are provided in `/test/test_data`. This technique allows to identify and correct issues in business logic more efficiently. Different test cases may be created by duplicating these files in a text editor;
- Setup an HTTP server to test input data provided as HTTP requests. Each dataset is available by a direct web URL;
- Copy the 3 datasets to HDFS and execute the test with the complete data;
- Setup different environments to verify performance;
- Execute a minimal of 10 tests per environment.

5.2 Environments

Document `/doc/execution.xlsx` presents the available environments and test results. The following environments are set:

- **Windows 10 standalone** - A Windows 10 machine where the IDE is installed, and the application is launched. This environment is also set for full debugging capacity;
- **Spark cluster in Linux** – With 4 virtual nodes and 8 cores total. It allows to test the application in real cluster mode, optimizing data serialization and shuffling;
- **iMac standalone** – An iMac from 2017, with considerably better HW than the other 2 environments.

5.3 Results

Document `/doc/execution_times.xlsx` gathers **10 executions** of Consultation application per environment. The statistics from the execution are resumed in the document **Error! Reference source not found.**, main results are:

- The slowest environment is Windows with an average total time ~ 30 minutes. This is explained by old HW and a complex environment of IDE and other tools running at the same time;
- The fastest environment is iMac with an average total time ~ 10 minutes. This is explained by fastest SSD disks, fastest memory and a better CPU;
- Linux cluster is closed to iMac values and performs very well in Consolidation creation, where most of the processing is done;
- Linux cluster loses to iMac in data parsing, this is due to much slower disks;
- Per environment, there are important differences between minimal and maximal execution times. This is explained by the distributive nature of Spark.

The final Consolidation dataset produced has a size of **20 024 474 x 10 and 1 708 723 881 bytes**. This means that **82%** of the initial Prescription entries are enriched with the necessary Physician and Drugs data. The gap is explained per **NPIs** reported in Prescription, but not available in the Physicians. No errors are detected in input data, besides the header elimination.

Logs, error dumps and samples of the first and last data 1000 entries in final **Consolidation.csv** dataset is available in `/test/<environment>/results/`. Data generation is the same for the 3 environments.